

University of Southern California

Viterbi School of Engineering

EE595

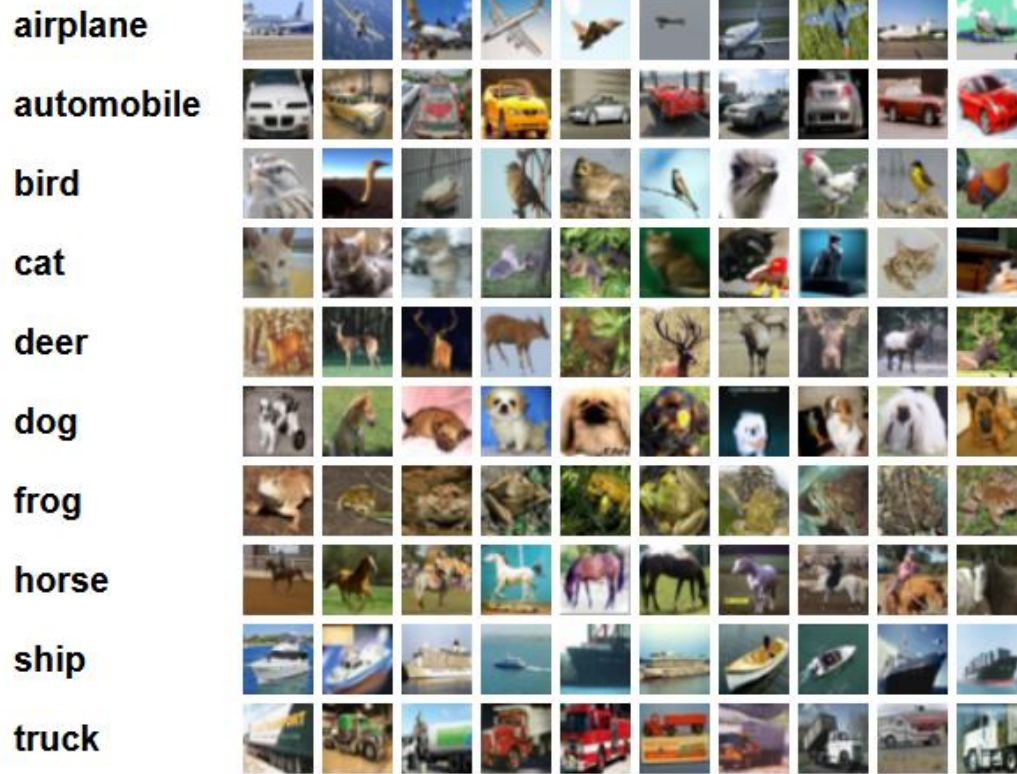
Lab 7 KNN, Naïve Bayes

Pengmiao Zhang

1. Image Classification with PCA and KNN

- 1. Dataset: CIFAR-10

Here are the classes in the dataset, as well as 10 random images from each:



Total 5 batches, total 50000 images for training and 10000 images for testing.

We use first 1000 images from the data batch 1

1. Image Classification with PCA and KNN

- 1. Dataset: CIFAR-10
- CIFAR-10 Website:

Dict Type

```
def unpickle(file):  
    import pickle  
    with open(file, 'rb') as fo:  
        dict = pickle.load(fo, encoding='bytes')  
    return dict
```

An image:

red 32*32
green 32*32
blue 32*32

•data -- a 10000x**3072** numpy array of uint8s. Each row of the array stores a **32x32** color image:

The first 1024 red channel values, the next 1024 the green, the final 1024 the blue.

The image is stored in **row-major** order, so that the first 32 entries of the array are the red channel values of the first row of the image.

•labels -- numbers in the range 0-9.

1. Image Classification with PCA and KNN

- **2. Dimension reduction:**
- **2.1 Convert the RGB images to grayscale with the following formula manually:**


$$Y = 0.299R + 0.587G + 0.114B$$

- **2.2 Use PCA to reduce dimensions.**
 - **Compute the PCA transformation by using only the training set with the sklearn PCA package**
 - **Perform dimensionality reduction on both training and testing sets to reduce the dimension of data from 1024 to D.**
 - **full SVD solver**


1. Image Classification with PCA and KNN

- **3. KNN Classifier from scratch:**
 - 3.1 use features obtained by PCA.
 - 3.2 use Manhattan distance
 - 3.3 K is specified as an input argument
 - 3.4 you are NOT allowed to use any library providing KNN classifier (e.g. sklearn)

$$d = \sum_{i=1}^n |x_i - y_i|$$



Manhattan



Euclidean

1. Image Classification with PCA and KNN

- **4. KNN Classifier with sklearn:**
- **Use sklearn package to do the KNN part again.**
- **Compare the results to verify your implementation.**

1. Image Classification with PCA and KNN

- **5. Program script and output**
- **5.1 script:**
- `python knn_<student_id>.py K D N PATH_TO_DATA`
- **5.2 output:**
 - predicted label and a ground truth label, separated by a single space
 - prediction accuracy
 - **two files**
 - `knn_results_sklearn.txt`
 - `knn_results.txt`

1	2
3	3
4	7
8	1
1	6
...	
...	
0.2	

2. Naïve Bayes

- 1. Dataset

UCI Seeds dataset
7 features, 3 types

15.26	14.84	0.871	5.763	3.312	2.221	5.22	1
14.88	14.57	0.8811	5.554	3.333	1.018	4.956	1
14.29	14.09	0.905	5.291	3.337	2.699	4.825	1
13.84	13.94	0.8955	5.324	3.379	2.259	4.805	1
16.14	14.99	0.9034	5.658	3.562	1.355	5.175	1
14.38	14.21	0.8951	5.386	3.312	2.462	4.956	1
14.69	14.49	0.8799	5.563	3.259	3.586	5.219	1

Attribute Information:

To construct the data, seven geometric parameters of wheat kernels were measured:

1. area A,
2. perimeter P,
3. compactness $C = 4 \cdot \pi \cdot A / P^2$,
4. length of kernel,
5. width of kernel,
6. asymmetry coefficient
7. length of kernel groove.

All of these parameters were real-valued continuous.

txt, space separation

pandas.read_table: dataframe

2. Naïve Bayes

- 2. Naïve Bayes from Scratch

$$\begin{aligned} h_{MAP} &\equiv \operatorname{argmax}_{h \in H} P(h \mid D) \\ &= \operatorname{argmax}_{h \in H} \frac{P(D \mid h)P(h)}{P(D)} \\ &= \operatorname{argmax}_{h \in H} P(D \mid h)P(h) \end{aligned}$$

2. Naïve Bayes

- **3. Algorithm**

- **1) log likelihood:** $\log p(\mathbf{X}|Y) = \sum_{i=1}^n \log p(X_i|Y)$

$$p(X_i = x_i | Y = y) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp - \frac{(x_i - \mu_i)^2}{\sigma^2}$$

conditional mean μ_i and variance σ_i^2 for each feature

- **2) Estimate prior distribution $p(Y)$**

- **3) In test, log-likelihood**

$$\hat{y} = \operatorname{argmax}_y \log p(\mathbf{X} = \mathbf{x}, Y = y)$$

2. Naïve Bayes

- **4. Comparison with the Implementation in Sklearn**

`sklearn.naive_bayes.GaussianNB`

- **5. Playing with Other Classifiers**

SVM, Decision Trees, Random Forests, Multi-layer Perceptron (MLP) classifier

2. Naïve Bayes

- **6. Script**

```
python naive_bayes_<STUDENT_ID>.py
```

- **5. Output**

No output generation, print information to console in this format:

My Naive Bayes:

Training acc: XX.XX% Training time: XXXX s

Testing acc: XX.XX% Testing time: XXXX s

Sklearn Naive Bayes:

Training acc: XX.XX% Training time: XXXX s

Testing acc: XX.XX% Testing time: XXXX s

<Other Classifier>:

Training acc: XX.XX% Training time: XXXX s

Testing acc: XX.XX% Testing time: XXXX s

3. Extra Credit

- **1. linear SVM**
 - Generate your own dataset randomly
 - LIBSVM
- **2. Implementing Multilayer Perceptron**
 - MNIST
 - pytorch